

Video Prompt Writing Guide

Base modes and full-reference mode, combined

Contents

1	Base Modes: T2VA, I2VA, FL2VA, L2VA	2
1.1	Task Overview	2
1.2	Final Prompt Structure	2
1.2.1	Part One Is the Instruction	2
1.2.2	Part Two Contains the Three Core Fields	3
1.3	Incorporating Keyframes into the Multimodal Description	3
1.3.1	I2VA: Begin from the Image and Develop Forward	3
1.3.2	FL2VA: Describe the Path Between the First and Last Frames	3
1.3.3	L2VA: Infer the Opening and Land on the Image at the End	3
1.4	Writing the Three Core Sections	4
1.4.1	Develop the Multimodal Description Along the Timeline	4
1.4.2	Shots and Cuts	4
1.4.3	Camera Motion: Motion Type, Amplitude, and Speed	4
1.4.4	Speakers, Dialogue, and Singing	5
1.4.5	On-Screen Text	6
1.4.6	overall_soundscape	6
1.4.7	non_diegetic_music	7
1.5	Base-Mode Cases	7
1.5.1	Case 1: T2VA	7
1.5.2	Case 2: I2VA	7
1.5.3	Case 3: FL2VA	8
1.5.4	Case 4: L2VA	8
2	Full-Reference Mode	9
2.1	Overall Structure	9
2.2	Reference Labels and Definitions	10
2.2.1	Subject Labels	10
2.2.2	Picture Labels	11
2.2.3	Video Labels	11
2.2.4	Audio Labels	11
2.2.5	Visual and Audio Tracks from the Same Reference Video	12
2.3	Writing the summary	12
2.4	Writing the retention_analysis	13
2.4.1	Visible Content	13
2.4.2	Audio	14
2.5	Writing the detailed_description	15
2.5.1	Basic Format	15
2.5.2	Differences from Base Mode	15
2.5.3	Using Reference Labels in Shots	16
2.5.4	Speakers, Audio Sources, and Dialogue	16

2.6	Sound Sections in Full-Reference Mode	17
2.7	Full-Reference Complete Example	18
3	Output Checklist	19

1 Base Modes: T2VA, I2VA, FL2VA, L2VA

1.1 Task Overview

Mode	Input	What the description must supply
T2VA	Text only	A complete audiovisual timeline built from the text
I2VA	First frame	T2VA body, a first-frame instruction, and a visual path developing forward from the first frame
FL2VA	First and last frame	T2VA body, a first-and-last-frame instruction, and a continuous path from the first frame to the last frame
L2VA	Last frame	T2VA body, a last-frame instruction, and a path converging from a plausible preceding state onto the last frame

1.2 Final Prompt Structure

1.2.1 Part One Is the Instruction

T2VA has no image-alignment instruction and begins directly with the three core fields.

I2VA always uses:

For the target video, at 0.00 seconds into the target video, <Picture 1> (from [Shot 1]) is fully referenced.

FL2VA always uses:

How the reference pictures align with the target video – Picture 1 (from Shot 1) aligns with the 0.00-second mark of the target video; Picture 2 (from Shot N) aligns with the S.SS-second mark of the target video.

L2VA always uses:

How the reference pictures align with the target video – <Picture 1> (from [Shot N]) aligns with the S.SS-second mark of the target video.

Here N is the index of the actual final shot, and S.SS is the effective video duration formatted to exactly two decimal places. The instruction must be the first line of the final prompt, followed by

one blank line before the core fields.

1.2.2 Part Two Contains the Three Core Fields

```
integrated_multimodal_description: [Shot 1] ...  
  
overall_soundscape: ...  
  
non_diegetic_music: ...
```

Field	Content
integrated_multimodal_description	Visuals, actions, shots, speakers, dialogue, singing, and diegetic audio along the timeline
overall_soundscape	Ambient sound, physical action sounds, and non-verbal human sounds across the entire video
non_diegetic_music	Background music the characters cannot hear and only the audience can hear

1.3 Incorporating Keyframes into the Multimodal Description

1.3.1 I2VA: Begin from the Image and Develop Forward

<Picture 1> is the actual first frame of the video at 0.00 seconds and belongs to [Shot 1]. The description should first establish the style, subjects, composition, and scene anchors in the image, then describe the next action. Character identity, clothing, colors, key objects, and spatial relationships should remain consistent.

Recommended structure: **first-frame anchor, action onset, continuous development, then result or reaction.**

1.3.2 FL2VA: Describe the Path Between the First and Last Frames

Picture 1 is the opening and Picture 2 is the ending. Focus on how the subject moves, how poses change, how objects are manipulated, how the composition evolves, and how the scene or lighting transitions.

FL2VA generally favors a single shot so the model can interpolate continuously from the first frame to the last frame. Use multiple shots only when they are explicitly specified. The last frame must be reached by the final [Shot N] at the end of the video.

Recommended structure: **first-frame state, observable intermediate changes, progressively narrowing differences, then last-frame state.**

1.3.3 L2VA: Infer the Opening and Land on the Image at the End

<Picture 1> is the final frame of the video and belongs to the last [Shot N]; it does not inherently belong to Shot 1. Infer a plausible earlier state from the user's intent and the last frame, then describe how the characters, objects, camera, and scene gradually approach the reference image.

Recommended structure: **plausible preceding state, explicit action and transition path, gradual convergence in the final shot, then last-frame landing.**

1.4 Writing the Three Core Sections

1.4.1 Develop the Multimodal Description Along the Timeline

`integrated_multimodal_description` is the main body of the rewritten prompt. Every detail should correspond to something visible or audible: visual style, initial composition, subject appearance and position, scene and key props, actions and reactions, shot changes, spoken language, and synchronized diegetic sound.

At the beginning of [Shot 1], state the overall style and initial composition. Common styles include Cinematic, live-action, 2D-animated, 3D CG, claymation, watercolor, and vintage film. For keyframe tasks, derive the style from the reference image; for T2VA, select it from the user's text.

```
[Shot 1] Live-action, cinematic, a medium-wide shot frames...
```

1.4.2 Shots and Cuts

Do not add a timestamp to the first shot. Use sequential shot numbers for later shots, and begin each one with a strictly increasing cut time that falls within the video duration:

```
[Shot 2] At 00:03.500, the camera cuts to...
```

For ordinary cuts, use the `camera cuts to`, the `shot cuts to`, the `shot transitions to`, the `shot changes to`, or the `shot switches to`. When explicitly requested by the user, cross-dissolve, fade, or wipe may also be used.

A cut should introduce new information about the subject, space, state, viewpoint, or time. If only the distance or a slight angle needs to change, prefer camera motion.

1.4.3 Camera Motion: Motion Type, Amplitude, and Speed

A complete camera-motion expression has three dimensions. The **motion type** defines how the camera moves, **amplitude** defines the range of compositional change, and **speed** defines the pacing of that change. Add amplitude and speed only when they are meaningful; medium amplitude and normal speed are usually omitted.

Motion type

Expression	Description
Zoom In / Zoom Out	The focal length changes while the camera body remains stationary
Push In / Pull Out	The camera moves forward or backward
Pan Left / Pan Right	The camera remains in place while the lens pivots horizontally

Expression	Description
Truck Left / Truck Right	The camera translates horizontally
Tilt Up / Tilt Down	The camera remains in place while the lens pivots vertically
Pedestal Up / Pedestal Down	The entire camera moves upward or downward
Arc Shot	The camera moves in an arc around the subject
Tracking Shot	The camera follows a moving subject
Static Shot	The camera position and lens remain still
Shake Slightly / Shake Strongly	Slight or strong camera shake
POV	The subject's point of view
Roll Clockwise / Roll Counterclockwise	The camera rolls around the lens axis

Amplitude and speed

Expression	Description
with small amplitude	Small-range change
with large amplitude	Large-range change
at slow speed	Slow movement
at fast speed	Fast movement

Camera motion should be written as a natural English action within the shot, rather than stacked as separate labels at the end of a sentence:

The camera pushes in with small amplitude at slow speed toward the folded letter in her hands.
The camera pans right with large amplitude at fast speed, revealing the open doorway.
The camera holds a static shot as the runner exits the frame.

1.4.4 Speakers, Dialogue, and Singing

Subjects who speak, sing, or produce an off-screen human voice use stable IDs such as (S1) and (S2). When multiple already-numbered speakers speak or sing together, use a compound ID such

as (S1,S2). A speaker keeps the same ID across shots; characters who never vocalize receive no speaker ID.

When a speaker first appears, provide enough information from the visual and audio context to establish a stable identity: character type, age, gender, whether the person is on-screen, pitch, timbre, speaking rate, or accent.

Place the speaker's identifying phrase, ID, action, and delivery **outside** <d>. Inside <d>, include only the language tag and the actual user-provided spoken content. Preserve every original word and punctuation mark verbatim; do not translate or rewrite them.

```
The young woman with a quiet, breathy voice (S1) says: <d>[English] I get off at  
the next station.</d>  
The two children (S1,S2) shout together, <d>[English] Wait for us!</d>
```

For voiceover, use the exact phrase says in an off-screen voiceover. Immediately after every voiceover <d> block, state that the corresponding on-screen character's lips remain closed:

```
The man (S1) says in an off-screen voiceover: <d>[English] I still remember that  
road.</d> while his lips remain completely closed.
```

When the same line of dialogue or lyrics crosses a cut, use <scenetrans> at the connecting points in both parts and explicitly state that the audio continues across the cut. Use <cutoff> when speech is truncated by the end of the video. Continuity may be expressed with continues seamlessly across the cut, continues uninterrupted into the next shot, carries over from the previous shot, or remains audible across the transition.

1.4.5 On-Screen Text

Place any banner, sign, label, subtitle, or neon text that is actually visible on screen in English double quotation marks. Preserve the original text and punctuation verbatim, without translation.

```
A red neon sign reading " 营业中 " glows above the doorway.
```

1.4.6 overall_soundscape

Use one to four English sentences in a single continuous paragraph to summarize the ambient sound, physical action sounds, and non-verbal human sounds across the full video: wind, rain, traffic, footsteps, fabric movement, impacts, breathing, laughter, or panting.

Dialogue, singing, and diegetic music already belong in the multimodal description and should not be repeated here. Use N/A only when the user explicitly requests complete silence throughout the video.

```
overall_soundscape: Steady rain taps against the café windows while low room  
ambience continues underneath. The entrance bell rings once, followed by wet  
footsteps and the soft scrape of a chair.
```

1.4.7 non_diegetic_music

Use one to three English sentences to describe background music that the characters cannot hear and only the audience can hear. Focus on instrumentation, speed, rhythm, and dynamic changes; do not use abstract mood words or explain the emotional function of the score.

Singing, instruments, radio, television, or phone music audible to the characters are diegetic events and should appear in the multimodal description. Use N/A when there is no non-diegetic music.

```
non_diegetic_music: Sparse piano notes at a slow tempo, joined by sustained low strings that gradually increase in volume before fading out.
```

1.5 Base-Mode Cases

1.5.1 Case 1: T2VA

With no reference image, construct the complete timeline directly from the text. You may add scene, character, action, and sound details that remain consistent with the user's intent.

```
integrated_multimodal_description: [Shot 1] Live-action, cinematic, a medium-wide shot frames a baker opening the shutters of a small street bakery before sunrise. The camera pushes in with small amplitude at slow speed as the middle-aged baker with a calm, slightly raspy voice (S1) places a fresh loaf on the wooden counter and says: <d>[English] First batch of the morning.</d> [Shot 2] At 00:05.000, the camera cuts to a close-up of steam rising from the sliced bread while the baker's final words carry over from the previous shot.
```

```
overall_soundscape: Wooden shutters scrape open over a quiet street as trays clink softly inside the bakery. The doorbell rings once, followed by light footsteps and the crisp sound of bread being sliced.
```

```
non_diegetic_music: A soft acoustic-guitar pattern at a moderate tempo, joined by sparse upright-bass notes and a gentle fade at the end.
```

1.5.2 Case 2: I2VA

Write the first-frame instruction first, then use the subject, composition, and scene in Picture 1 as the starting point of Shot 1 before describing how the scene continues to develop.

```
For the target video, at 0.00 seconds into the target video, <Picture 1> (from [Shot 1]) is fully referenced.
```

integrated_multimodal_description: [Shot 1] Live-action, cinematic, the young woman shown in <Picture 1> remains beside the rain-covered train window, preserving her appearance, clothing, seat position, and the carriage layout. The camera trucks right with small amplitude at slow speed as she lifts her gaze from the folded letter toward the passing city lights. Her reflection moves across the glass while the quiet, breathy young woman (S1) says: <d>[English] I get off at the next station.</d> She folds the letter along its existing crease.

overall_soundscape: The train wheels produce a steady metallic rhythm beneath a low ventilation hum. Rain ticks against the window while paper rustles softly in her hands.

non_diegetic_music: Sustained cello notes at a slow tempo with widely spaced piano tones, gradually decreasing in volume.

1.5.3 Case 3: FL2VA

The two images anchor the opening and ending respectively. The body should not repeat two static image descriptions; instead, it should supply the motion path that connects them. The following example is an eight-second single shot.

How the reference pictures align with the target video – Picture 1 (from Shot 1) aligns with the 0.00-second mark of the target video; Picture 2 (from Shot 1) aligns with the 8.00-second mark of the target video.

integrated_multimodal_description: [Shot 1] Live-action, cinematic, a rain-soaked cyclist begins in the position and framing established by Picture 1, holding a closed black umbrella beside a silver bicycle. The camera pulls out with small amplitude at slow speed as she releases the bicycle handle, raises the umbrella above her shoulder, and presses the runner upward until the canopy opens. Water rolls from the expanding fabric while she steps beneath it, rotates the handle into the final angle, and settles into the pose, spacing, and composition established by Picture 2 at the end of the shot.

overall_soundscape: Rain falls steadily on the pavement, followed by the metallic click of the umbrella runner and the soft snap of the canopy opening. Water drips from the bicycle frame as distant traffic passes.

non_diegetic_music: N/A

1.5.4 Case 4: L2VA

The image anchors only the final moment. First establish a compatible earlier state, then let the actions, object states, and composition gradually land on Picture 1 in the final shot. The following example is a six-second single shot.

How the reference pictures align with the target video – <Picture 1> (from [Shot 1]) aligns with the 6.00-second mark of the target video.

`integrated_multimodal_description`: [Shot 1] Live-action, cinematic, a close shot begins with an intact drinking glass near the edge of a dark wooden table, while the same hand and sleeve visible in <Picture 1> approach from the right. The camera pushes in with small amplitude at slow speed as the fingertips strike the rim. The glass tips, falls, and hits the floor with a sharp impact; cracks spread through it as fragments slide outward. Toward the end, the moving pieces lose momentum and settle into the exact broken arrangement, hand position, camera angle, lighting, and final composition established by <Picture 1>.

`overall_soundscape`: Fingertips tap the glass before it scrapes across the tabletop, falls, and breaks with a sharp crash. Small fragments scatter and gradually stop sliding across the floor.

`non_diegetic_music`: A low electronic pulse at a slow tempo, ending immediately after the glass breaks.

2 Full-Reference Mode

This part explains how rewrite outputs are organized and written in full-reference mode. The basic formats for shots, camera movement, speakers, dialogue, and ordinary sound are shared with Section 1; this part focuses on the reference labels, analysis sections, and format differences specific to full-reference mode.

Write all six rewrite sections in English. Preserve the original language only for dialogue and lyrics inside <d> and for text visibly present in the scene.

Description detail. Make `detailed_description` as detailed and explicit as possible. For each shot, clearly establish the current composition, subject appearance and position, environment and lighting, actions and state changes, camera movement, current sound, and the points where referenced content actually appears or takes effect. Avoid reducing the description to a plot summary or a list of reference relationships.

2.1 Overall Structure

A complete rewrite output consists of six sections in the following order:

Section	Purpose
<code>subject_definitions</code>	Defines referenced content and its reference labels
<code>summary</code>	Summarizes the task type, target video, and main reference relationships

Section	Purpose
<code>retention_analysis</code>	Describes how referenced content is preserved, transferred, or reused
<code>detailed_description</code>	Describes visuals, actions, shots, sound, and dialogue in playback order
<code>overall_soundscape</code>	Summarizes ambience and physical sounds
<code>non_diegetic_music</code>	Describes background music audible only to the audience

2.2 Reference Labels and Definitions

Full-reference rewrites use four types of labels to identify the source and role of referenced content:

Label	Meaning
<code><Subject N></code>	Visible content abstracted from reference assets that can be reused or modified in the target video
<code><Picture N></code>	A reference image used as a concrete target frame or shot-planning anchor
<code><Video N></code>	A reference video that provides an editing source, continuation starting point, or whole-video temporal structure
<code><Audio N></code>	An audio signal that is copied or referenced

Once a reference label is assigned to a piece of content, it keeps the same meaning across `subject_definitions`, `summary`, `retention_analysis`, `detailed_description`, and the audio sections.

`subject_definitions` defines each piece of referenced content that must be tracked separately later, such as a person, an environment, a source video's structure, or an audio track. Give each item its own line and explain what its label denotes, its reference role, and the main features to follow; name the corresponding source asset when its provenance needs to be made explicit. If `<Picture N>` or `<Video N>` only identifies the source of another referenced item and will not be analyzed or used separately later, cite it inside that item's definition without adding a separate line.

2.2.1 Subject Labels

`<Subject N>` is used for reusable visible content, including:

- People, animals, or objects
- Scenes, backgrounds, or environments
- Clothing, props, interfaces, or visual effects
- Styles, actions, expressions, or poses

It represents a content unit that will actually be used in the target video, rather than the source file itself. One subject may be defined by multiple reference assets, and one reference asset may provide multiple subjects.

```
<Subject 1> is the young woman in <Picture 1>, with long dark hair, a blue cardigan, and a thin silver necklace.
```

When the same subject comes from multiple assets, combine the sources and state what each asset provides:

```
<Subject 1> is the woman whose appearance comes from <Picture 1> and whose walking motion comes from <Video 1>.
```

2.2.2 Picture Labels

Use a standalone <Picture N> when the reference image itself serves as a shot's first frame, keyframe, last frame, edited keyframe, or composition anchor:

```
<Picture 2> is the first frame of [Shot 1], showing a woman seated beside a café window.
```

If an image is used only to define a character, scene, costume, or style, do not create a standalone picture entry. Instead, cite the image source inside the corresponding <Subject N> definition.

When an image acts as a storyboard or shot-planning reference, state which shots it maps to and what planning information it provides:

```
<Picture 3> is a storyboard reference for [Shot 1] and [Shot 2], defining their viewpoint, subject placement, and shot order.
```

2.2.3 Video Labels

<Video N> is reserved for whole-video relationships, such as:

- Editing an original video
- Continuing from the end of an original video
- Referencing the original video's camera movement, cuts, rhythm, or temporal structure

```
<Video 1> is the source video for the target video edit.
```

If a person, object, scene, action, or effect from a reference video is reused as visible content, it still belongs under <Subject N>. <Video N> identifies the asset or structural source and does not replace subject labels.

2.2.4 Audio Labels

<Audio N> represents a standalone audio asset or an enabled synchronized audio track from a reference video. Common uses include:

- Copying all or part of an audio signal
- Referencing a background-music style
- Referencing a speaker's voice timbre and delivery
- Using dialogue, lyrics, or sound effects from the original audio
- Referencing beat, rhythm, or audio continuity

When an <Audio N> explicitly corresponds to a target speaker, reuse that speaker's global ID in the definition: write <Subject N> (Sx) when the speaker maps to a defined subject, or use a stable voice description followed by (Sx) otherwise. The ID comes from the target video's global speaker order and is not independently assigned or renumbered in the audio definition.

```
<Audio 1> is the voice-timbre reference for <Subject 1> (S1).
```

When one audio asset serves multiple roles, describe those roles in one natural sentence rather than creating additional subsections.

2.2.5 Visual and Audio Tracks from the Same Reference Video

<Video N> and <Audio N> are numbered independently. Each index indicates only the label's order within its own category and does not encode a pairing between the two categories. The same reference video may therefore correspond to <Video 1> and <Audio 2>; different indices do not prevent them from coming from the same source asset.

An ordinary reference video does not create <Audio N> merely because the file contains sound.

An <Audio N> definition primarily states the audio's role and does not have to name the <Video N> it comes from. State the shared source only when needed to remove provenance ambiguity:

```
<Video 1> is the source video for the target video edit.  
<Audio 2> is the synchronized audio track of <Video 1> and is reused in the target video.
```

2.3 Writing the summary

This section uses one short English paragraph to summarize the target video and its reference relationships. It begins with a square-bracketed task-type prefix:

```
[reference generation] ...  
[video editing + reference generation + audio reuse] ...
```

Choose task types according to the actual role each reference asset plays in the target video:

Task type	When to use it
keyframe completion	An image serves as the target video's first frame, keyframe, last frame, edited keyframe, or another concrete frame anchor
reference generation	An image, video, or audio asset provides generation guidance for a character, scene, style, action, camera movement, storyboard, and so on, without serving as a concrete frame or as the source video being edited or continued
video editing	An existing source video is directly modified; editing an image or generating between still keyframes does not belong to this type
video continuation	New content continues, extends, resumes, or transitions from an existing source video
audio reuse	The same audio signal is reused in full or in part
audio reference	The audio signal is not copied directly; only its music style, timbre, dialogue or lyric content, sound-effect texture, beat, or continuity is referenced

When a task satisfies multiple relationships, combine the task types with + and do not repeat a type. For example, continuing from a source video while using an image as the last frame is written as [video continuation + keyframe completion]. Editing a source video while retaining its original audio may be written as [video editing + audio reuse].

The mere presence of video or audio does not automatically create a corresponding task type. If a reference video provides only camera movement, cuts, or rhythm, it normally belongs to reference generation. Use video editing or video continuation only when that video is directly edited or continued.

When editing a source video, use audio reuse as well if its original audio remains audible. When continuing a source video without directly copying the audio signal, use audio reference if the new audio only continues the original track's audible characteristics.

The summary uses the previously defined labels to describe the main subjects, shot flow, and roles of the reference assets. Do not introduce new reference labels in this section.

For video-editing tasks, begin the summary after the task-type prefix with:

```
The target video is an edited version of <Video 1>.
```

2.4 Writing the retention_analysis

This section describes how each piece of referenced content is preserved, transferred, copied, or referenced in the target video. Use one line for each reference label and preserve the meaning established in subject_definitions.

2.4.1 Visible Content

<Subject N>, <Picture N>, and <Video N> use the following relationship markers. These markers are fixed English values in the output format:

Relationship marker	Meaning
fully_preserved	The defined role of the referenced content is fully preserved
partially_preserved	The referenced content is still used, but some defined characteristics are changed or only partially retained
attribute_transfer	Referenced characteristics are transferred to a different identifiable target subject
weak_reference	Only broad similarity in style, category, composition, or atmosphere is retained

```
<Subject 1> (appears in [Shot 1], [Shot 3]): fully_preserved - ...
<Picture 2> ([Shot 1] first frame): fully_preserved - ...
<Video 1> (cut and pacing structure): weak_reference - ...
```

2.4.2 Audio

<Audio N> uses the following relationship markers:

Relationship marker	Meaning
fully_copy	The complete source audio serves as the target video's complete final audio track
partially_copy	Only part of the timeline or selected audio layers are copied, or other sounds are added, removed, or replaced after copying
reference	The signal is not copied directly; only timbre, rhythm, music style, dialogue content, or sound texture is referenced
weak_reference	Only broad similarity in category or atmosphere is retained

```
<Audio 1>: fully_copy - <Audio 1> is reused 1:1 as the target video's complete
final audio track.
<Audio 2>: reference - the target speaker follows <Audio 2>'s voice timbre and
measured delivery without copying the original signal.
```

Choose each relationship marker only within the reference role already defined for that label in `subject_definitions`. Do not treat newly added actions, backgrounds, or plot events in the target video as losses of reference fidelity.

2.5 Writing the detailed_description

This is the main body of a full-reference rewrite. It describes visuals, actions, sound, and dialogue shot by shot in target-video playback order and inserts reference labels where they apply.

2.5.1 Basic Format

The basic format follows Section 1:

- Write the body in English. Preserve the original language of dialogue, lyrics, and visible text.
- `[Shot 1]` marks the opening shot and has no timestamp. Later shots use `[Shot N]` At `MM:SS.mmm`, ... to mark cut times.
- Write camera movement as natural English within the current shot, including movement type, amplitude, and speed when they need to be expressed.
- Give vocal sources stable (S1), (S2), and subsequent IDs. Write dialogue and lyrics as `<d>[Language] ...</d>`.
- Use `<scenetrans>`, `<cutoff>`, and the corresponding continuity descriptions for dialogue crossing a cut, speech truncated by the video ending, and continuous audio across shots.

For complete rules and examples covering camera vocabulary, group speech, voice-over, dialogue across cuts, and visible text, see Section 1.

2.5.2 Differences from Base Mode

Dimension	Base modes	Full-reference mode
Main field	<code>integrated_multimodal_description</code>	<code>detailed_description</code>
Style opening	Written after <code>[Shot 1]</code>	Established in one or two English sentences before <code>[Shot 1]</code>
Reference information	Does not use full-reference labels	Inserts <code><Subject N></code> , <code><Picture N></code> , <code><Video N></code> , and <code><Audio N></code> at their first appearance and where their roles apply
Audio relationships	Describes the target video's own sound	Cites <code><Audio N></code> in the corresponding shot or audio phase and states whether the signal is copied or referenced

Opening example:

The target video is in a cinematic, literary music-video style with soft lighting and a slightly desaturated color palette.
[Shot 1] The scene opens in a crowded urban street...

[Shot 2] At 00:09.000, the shot cuts to an extreme close-up...

For generation tasks, `detailed_description` is normally 350 to 500 English words. Dialogue-dense content prioritizes fitting the complete spoken timeline rather than mechanically reaching a word count. Video-editing descriptions scale with the complexity of the source video and do not have to follow the generation-task range. A single shot does not automatically justify a shorter description; distribute detail across multiple shots according to their information load.

2.5.3 Using Reference Labels in Shots

At the first clear appearance of an important `<Subject N>`, describe its referenced characteristics, position in the frame, and current action within what is actually visible in the shot. Continue using the same label in later shots without redefining what the label represents.

Use natural phrasing for concrete frame anchors:

```
the shot begins from <Picture 1>
the shot's keyframe corresponds to <Picture 2>
the shot ends on <Picture 3>
```

When editing or continuing an original video, cite `<Video N>` naturally where its source state, structure, or continuation relationship applies. Cite `<Audio N>` in the shot or semantic phase where the audio relationship is active.

2.5.4 Speakers, Audio Sources, and Dialogue

The basic speaker-ID and `<d>` formats follow Section 1. When a referenced subject physically speaks, retain both the visual reference label and the speaker ID:

```
<Subject 2> (S1) turns toward the woman and says, <d>[English] Last summer, I went
to my grandfather's house.</d>
```

`<Subject N>` identifies the referenced subject, while `(Sx)` identifies the actual speaker. When the subject speaks, write `<Subject N> (Sx)`. If the same subject speaks off-screen, keep the same form and mark it as off-screen. When the speaker does not correspond to a defined subject, use a stable voice description followed by `(Sx)`.

When verbal content is only a cue within a directly reused BGM or complete soundtrack, and no person, character, narrator, or other independent vocal source physically produces it, use `<Audio N>` as the audible source and do not invent an additional `(Sx)`. If a concrete person, character, narrator, or other independent vocal source produces the voice, assign and reuse `(Sx)` for that source:

```
When <Audio 1> reaches the phrase <d>[English] I'm lonely lonely lonely I'm
lonely</d>, <Subject 1> performs the corresponding hand gesture without becoming a
separate speaker source.
```

When dialogue, narration, or lyrics from reference audio are directly reused, or when the input prompt explicitly requests their reperformance, preserve the exact source words and original lan-

guage inside `<d>`. Write `[unclear]` for unintelligible spans instead of guessing or paraphrasing them. Standardize punctuation to the basic written marks needed to express the sentence, such as `,`, `.`, `?`, and `!`; remove repeated tildes, emoji, bullets, and repeated or decorative punctuation. End complete statements, questions, and exclamations with `.`, `?`, or `!` respectively before `</d>`.

When only timbre, rhythm, emotion, or delivery is referenced, do not carry the original dialogue from the reference audio into the target video.

Assign (Sx) once according to the order of actual vocal events in the target video. Reuse the corresponding ID at every actual vocal event in `detailed_description`; an `<Audio N>` definition bound to a target speaker in `subject_definitions` also reuses the same (Sx) but never assigns a new one independently. Do not write (Sx) in `retention_analysis`.

2.6 Sound Sections in Full-Reference Mode

The definitions of these two sound categories follow Section 1.

`overall_soundscape` summarizes ambience and physical sounds across the full video. Dialogue, singing, and sound events synchronized to a particular shot remain in `detailed_description`:

```
overall_soundscape: Quiet indoor room tone and a low ventilation hum continue throughout the video.
```

`non_diegetic_music` describes background music that the characters cannot hear and that is audible only to the audience. When music is present, state its instrumentation, tempo, and dynamic development:

```
non_diegetic_music: A restrained solo-piano score at a slow tempo, with sustained low cello underneath and no swell.
```

When reference audio is used, state its copy or reference relationship only in the section that matches the audible layer: ambience and sound effects belong in `overall_soundscape`, while audience-only score belongs in `non_diegetic_music`. If the same audio provides both kinds of content, describe the corresponding relationship in each section:

```
overall_soundscape: The copied ambience layer from <Audio 1> continues throughout the target video.  
non_diegetic_music: <Audio 2> is directly reused as the complete audience-only score.
```

Write complete dialogue and lyrics only inside `<d>` in `detailed_description`; do not repeat them in these two sections.

2.7 Full-Reference Complete Example

subject_definitions:

<Subject 1> is the coffee-shop environment in <Picture 1>, featuring an exposed brick wall, an orange tufted sofa with patterned pillows, a neon sign, and a wooden coffee table.

<Subject 2> is the fluffy white Samoyed in <Picture 2>, <Picture 3>, and <Picture 4>, with thick white fur, pointed ears, a dark nose, and a curved tail.

<Subject 3> is the young blonde woman in <Video 1>, with long blonde hair and a light-pink button-down shirt with rolled-up sleeves.

<Subject 4> is the young man in <Video 2>, with short wavy brown hair and a dark-grey hoodie with drawstrings.

<Audio 1> is the voice-timbre reference for <Subject 3> (S1), containing a spoken English vocal layer.

summary:

[reference generation + audio reference] The target video shows <Subject 3> eating a cookie in <Subject 1>. <Subject 4> enters with <Subject 2>, which lunges toward the cookie. The three-shot exchange uses <Audio 1> as the voice-timbre reference for <Subject 3> and ends with a canned audience laugh.

retention_analysis:

<Subject 1> (appears in [Shot 1], [Shot 2], [Shot 3]): fully_preserved - the exposed brick wall, orange tufted sofa, patterned pillows, neon sign, and wooden coffee table are retained.

<Subject 2> (appears in [Shot 1], [Shot 2]): fully_preserved - the Samoyed's thick white fur, pointed ears, dark nose, and curved tail are retained.

<Subject 3> (appears in [Shot 1], [Shot 2], [Shot 3]): fully_preserved - the blonde woman's identity, long hair, and light-pink shirt are retained.

<Subject 4> (appears in [Shot 1], [Shot 2]): fully_preserved - the young man's short wavy brown hair and dark-grey hoodie are retained.

<Audio 1>: reference - its vocal timbre guides the dialogue delivery of <Subject 3> without copying the original signal.

detailed_description:

The target video uses a realistic multi-camera sitcom style with warm indoor lighting.

[Shot 1] A medium shot establishes <Subject 1>, the coffee shop with its exposed brick wall, orange tufted sofa, patterned pillows, neon sign, and wooden coffee table. <Subject 3> (S1), the young woman with long blonde hair and a light-pink button-down shirt with rolled-up sleeves, sits on the sofa holding a chocolate-chip cookie. From the left, <Subject 4>, the young man with short wavy brown hair and a dark-grey hoodie with drawstrings, enters holding the leash of <Subject 2>, the thick-furred white Samoyed with pointed ears, a dark nose, and a curved tail. The dog lunges toward the cookie and pulls the leash taut. <Subject 3> (S1) jerks her hand back and, using the clear youthful voice timbre referenced from <Audio 1>, exclaims with light annoyance, <d>[English] Hey! Watch your dog!</d> She closes her lips and guards the cookie while <Subject 4> pulls the dog back.

[Shot 2] At 00:03.000, the shot cuts to a close-up of <Subject 4> (S2), the young man in the dark-grey hoodie from Shot 1, sitting beside <Subject 3> on the sofa and holding <Subject 2> securely in his arms. <Subject 4> (S2) says in a casual young male voice with a playful tone and an easy conversational pace, <d>[English] He just likes cookies more than me.</d> He closes his mouth into an apologetic smile and strokes the dog's thick white fur.

[Shot 3] At 00:05.000, the shot cuts to a close-up of <Subject 3> (S1), the blonde woman in the light-pink shirt from Shot 1. Her annoyance softens as she looks toward the Samoyed. <Subject 3> (S1) replies in the same clear youthful voice referenced from <Audio 1> with an amused cadence, <d>[English] Well, he has good taste at least.</d> She smiles and raises the cookie in a small toast-like gesture. A classic canned audience laugh begins immediately after the line and continues through the final frame.

overall_soundscape:

Soft indoor coffee-shop room tone continues throughout the scene.

non_diegetic_music:

N/A

3 Output Checklist

All modes

- Correct instruction line present as the first line, followed by a blank line, or absent for T2VA
- S.SS duration formatted to exactly two decimal places, and N is the real final shot index
- [Shot 1] carries no timestamp; later cut times strictly increase and fall inside the duration
- Every cut introduces new subject, space, state, viewpoint, or time information
- Camera motion written as natural action, with amplitude and speed only where meaningful
- Speaker IDs stable across shots; characters who never vocalize carry no ID
- <d> contains only the language tag and verbatim speech
- Voiceover uses says in an off-screen voiceover and a closed-lips statement
- <scenetrans> present in both parts when a line crosses a cut; <cutoff> used when speech is truncated

- On-screen text in double quotation marks, untranslated
- `overall_soundscape` is one to four sentences with no dialogue, singing, or diegetic music
- `non_diegetic_music` is one to three sentences covering instrumentation, tempo, and dynamics
- N/A used only where genuinely warranted

Full-reference mode

- All six sections present and in order
- Every label defined once and used consistently across all sections
- Task-type prefix matches the actual role of each reference asset
- No new reference labels introduced in `summary`
- No (Sx) anywhere in `retention_analysis`
- Relationship markers chosen within each label's defined role
- Style established in one or two sentences before `[Shot 1]`